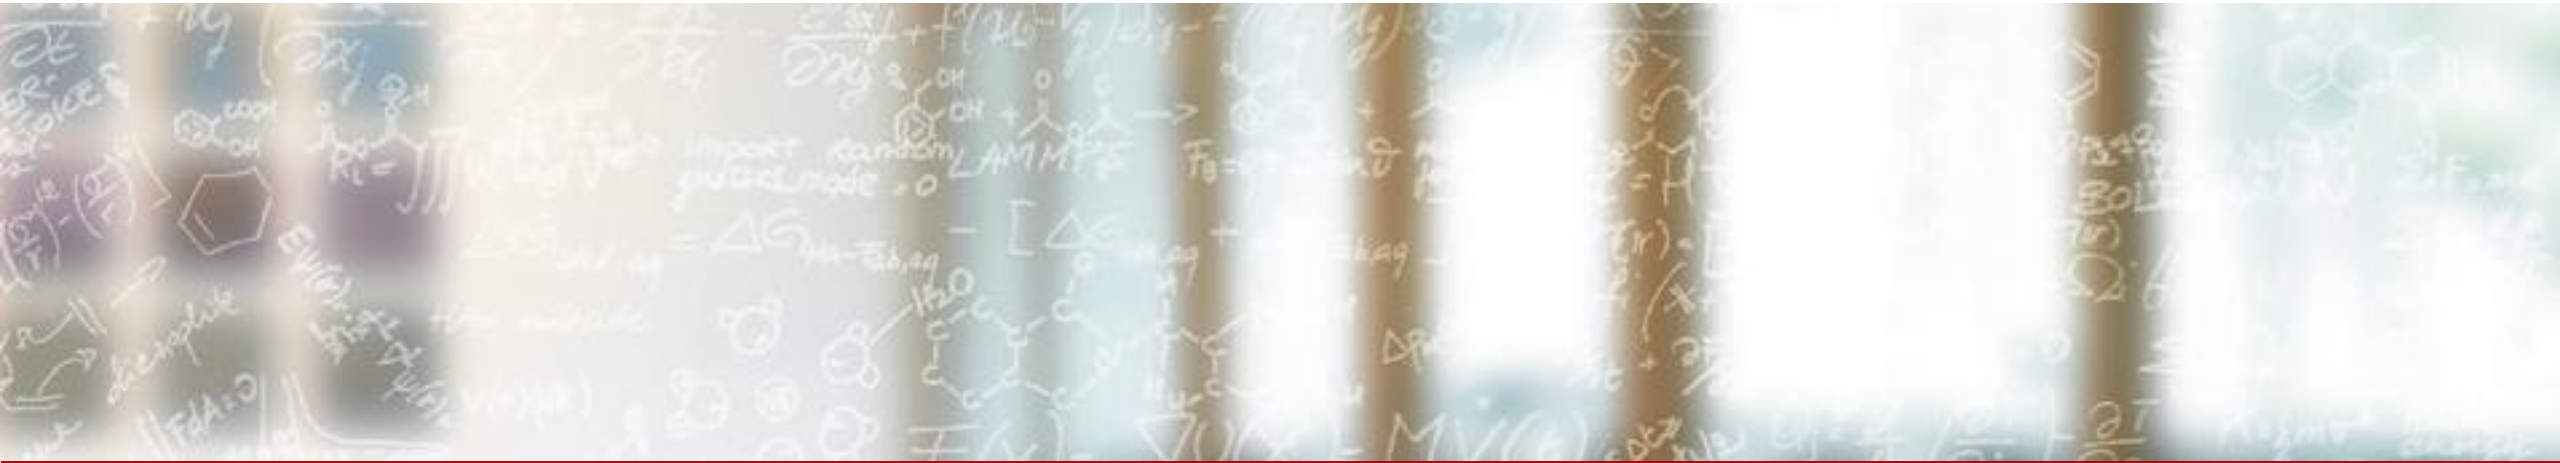




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Containerized CI/CD on Alps

Theofilos Manitaras, CSCS

Andreas Fink, CSCS

October 4, 2023

# Table of Contents

1. Creating container images
2. Testing container images
3. Embracing multi-stage builds & image splitting
4. Building for the target architecture
5. Image labels

# Creating container images

---

# Creating container images

Creating container images is straightforward:

```
include:  
  - remote: 'https://gitlab.com/cscs-ci/recipes/-/raw/master/templates/v2/.ci-ext.yml'  
  
stages:  
  - build  
  
build software:  
  extends: .container-builder  
  stage: build  
  variables:  
    PERSIST_IMAGE_NAME: $CSCS_REGISTRY_PATH/hello_world:01-$CI_COMMIT_SHORT_SHA  
    DOCKERFILE: ci/docker/Dockerfile
```

 **Needed for Image Building**

 **Mandatory variables**

You can find extra information and options regarding building container images [here](#)

# Testing container images

---

# Testing container images

After building a container image, it's essential to test/run it on the target system(s):

```
.  
.br/>test software:  
  extends: .container-runner-daint-mc  
  stage: test  
  image: $CSCS_REGISTRY_PATH/hello_world:01-$CI_COMMIT_SHORT_SHA  
  script:  
    - /build/hello  
  variables:  
    SLURM_NTASKS: 4  
    SLURM_JOB_NUM_NODES: 1  
    SLURM_LABELIO: 1  
    USE_MPI: 'YES'
```

**Selecting the corresponding runner**

**Container command to execute**

**Variables controlling slurm allocation & Sarus parameters**

You can find extra information and options regarding running containers [here](#)



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Embracing multi-stage builds & reusing images

---

# Multi-stage builds

Multi-stage builds are an essential mechanism during the image building:

- Multiple FROM statements are used in a Dockerfile
- Artifacts can be copied from one stage to another
- Allow creation of lightweight images containing only the essentials to run

```
# Start from "devel" image containing all the necessary components for building
FROM docker.io/nvidia/cuda:11.2.2-devel-ubuntu20.04 as builder
.
.
# Build binary <my_binary>
.
.
# Build the lightweight runtime image
FROM docker.io/nvidia/cuda:11.2.2-runtime-ubuntu20.04

COPY --from=builder <source path to binary> <target path>
```



# Image splitting (1/2)

A typical pattern emerging when using container images is the following:

- Image components e.g third party libraries, change at a much slower pace than the actual software being developed
- Building an image with all the components from scratch might take a long time, wasting computing resources

A solution to the above is to split the image in two:

1. An image containing the build environment, that can be updated on demand
2. An image that only builds the software to be tested/shipped, using the base image as it's starting point

## Image splitting (2/2)

An example pipeline is given [here](#):

```
build base:
  extends: .container-builder-dynamic-name
  stage: build_base
  variables:
    WATCH_FILECHANGES: ci/docker/Dockerfile.base
    DOCKERFILE: ci/docker/Dockerfile.base
    PERSIST_IMAGE_NAME: $CSCS_REGISTRY_PATH/base/hello_world_base

build software:
  extends: .container-builder
  stage: build
  variables:
    PERSIST_IMAGE_NAME: $CSCS_REGISTRY_PATH/hello_world:03-$CI_COMMIT_SHORT_SHA
    DOCKERFILE: ci/docker/Dockerfile
    DOCKER_BUILD_ARGS: '["BASEIMG=$BASE_IMAGE"]'
```

**Use to propagate BASE\_IMAGE**

**Build when base Dockerfile changes**

**Pass BASE\_IMAGE as build argument**

# Building for the target architecture

---

# Building for the target architecture

Container images should be built with the target running system in mind:

- Crucial in order to generate binaries compatible with system specific futures (e.g accelerators, high speed interconnects, vector instructions)
- Relying on autodetection of system characteristics should be avoided since, the system where the image is built might be different from the target one
- CSCS offers a growing list of essential building blocks to facilitate the compatibility and performance of the resulting images:
  - [Base images](#)
  - [CI pipeline templates](#)
  - Helper scripts (e.g spack-install-helper)

# Building for the target architecture example (1/2)

Different images can be built from scratch as follows:

```
FROM ubuntu:22.04 as builder
```

```
ARG TARGET=haswell
```



**Accept the architecture as a build argument**

```
.
```

```
.
```

```
.
```

```
RUN mkdir -p /opt/spack/spack-env/ \
```

```
.
```

```
.
```

```
&& (echo "spack: " \
```

```
&& echo "  packages: "\
```

```
&& echo "    all: "\
```

```
&& echo "      target: [${TARGET}]" \
```



**Use the architecture with Spack**

```
&& echo "  specs: " \
```

```
&& echo "    - osu-micro-benchmarks@6 ^${MPI_SPEC}" \
```

```
.
```


```
.
```

```
.
```

## Building for the target architecture example (2/2)

The corresponding pipeline looks as follows:


```
build haswell:
  extends: .container-builder
  stage: build
  variables:
    PERSIST_IMAGE_NAME: $CSCS_REGISTRY_PATH/mpich_osu:3.4.3_6.0_haswell
    DOCKERFILE: ci/mpi_osu/Dockerfile
    DOCKER_BUILD_ARGS: '["TARGET=haswell"]'
```



**Build for haswell**

```
build zen2:
  extends: .container-builder
  stage: build
  variables:
    PERSIST_IMAGE_NAME: $CSCS_REGISTRY_PATH/mvapich2_osu:2_6.0_zen2
    DOCKERFILE: ci/mpi_osu/Dockerfile
    DOCKER_BUILD_ARGS: '["TARGET=zen2"]'
```




**Build for zen2**




# Build/Run pipeline

## Final pipeline for different architectures

✓ Passed

CSCS Ciext created pipeline for commit `19156dc5`  finished 4 hours ago

For `--CSCSCI__pr4`

latest  4 Jobs  0  34 seconds, queued for 2 seconds

Pipeline

Needs

Jobs 4

Tests 0

build

run

✓ build haswell

✓ build zen2

✓ run haswell

✓ run zen2

# Image labels

---



# Labels in a Nutshell

Labels are a mechanism for adding useful metadata to a container image:

```
FROM docker.io/nvidia/cuda:11.2.2-devel-ubuntu20.04 as builder

# Labels for builder image
LABEL "rfm.features"="osu-micro-benchmarks;mpi;serial;openmp;mvapich2"
LABEL "rfm.cc"="mpicc"
LABEL "rfm.cxx"="mpic++"
LABEL "rfm.ftn"="mpif90"
.
.
.
# Labels for runner image
FROM docker.io/nvidia/cuda:11.2.2-runtime-ubuntu20.04

LABEL "maintainer"="cscs-ci-demo"
LABEL "rfm.features"="osu-micro-benchmarks;mpi;serial;openmp;cuda"
LABEL "rfm.cc"="mpicc"
LABEL "rfm.cxx"="mpic++"
LABEL "rfm.ftn"="mpif90"
```

# Accessing Labels

Labels can be easily accessed via image inspection, using [Docker](#), [skopeo](#), etc.

```
# Inspect using Docker
$ docker inspect <image_name> --format="{{json .Config.Labels}}" | jq
{
  "maintainer": "cscs-ci-demo",
  "rfm.cc": "mpicc",
  "rfm.cxx": "mpic++",
  "rfm.features": "osu-micro-benchmarks;mpi;serial;openmp;cuda",
  "rfm.ftn": "mpif90"
}


# Inspect using Skopeo
$ skopeo inspect docker://<image_name> | jq '.Labels'
{
  "maintainer": "cscs-ci-demo",
  "rfm.cc": "mpicc",
  "rfm.cxx": "mpic++",
  "rfm.features": "osu-micro-benchmarks;mpi;serial;openmp;cuda",
  "rfm.ftn": "mpif90"
}
```



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your attention.**